

14

Unit Testing with the Unit Test Framework

Unit testing involves writing code to verify a system at a lower and more granular level than with other types of testing. It is used *by* programmers *for* programmers and is quickly becoming standard practice at many organizations. The Developer and Tester editions of Visual Studio Team System now include unit testing features that are fully integrated with the IDE and with other features such as reporting and source control. Developers no longer need to rely on third-party utilities, such as NUnit, to perform their unit testing, although they still have the option to use them.

In this chapter, we describe the concepts behind unit testing, why it is important, and how to create effective unit test suites. We introduce the practice of *test-driven development*, which involves the creation of tests before writing the code to be tested.

We introduce the syntax of writing unit tests, and you will learn how to work with Team System's integrated features for executing and analyzing those tests. We then go into more detail about the classes available to you when writing your unit tests, including the core `Assert` class and many important attributes. You will also see how easy it is to create data-driven unit tests, whereby a unit test is executed once per record in a data source, and has full access to that bound data.

We also describe the features Team System offers for easily accessing nonpublic members from your unit tests. In addition, you will learn how Team System enables the generation of unit tests from existing code as well as the generation of member structures when writing unit tests.

Team System also provides built-in features to support the unit testing of ASP.NET applications. You will learn how to create ASP.NET unit tests that have full access to the ASP.NET run-time context. This enables actions such as setting and reading form control values and execution of those controls' events.

Finally, we describe how Team System can help measure the effectiveness of your tests. *Code coverage* is the concept of observing which lines of code are used during execution of unit tests. You can easily identify regions of your code that are not being executed and create tests to verify that code. This will improve your ability to find problems before the users of your system do.

Unit Testing Concepts

You've likely encountered a number of traditional forms of testing. Your quality assurance staff may run automated or manual tests to validate behavior and appearance. Load tests may be run to establish that performance metrics are acceptable. Your product group might run user acceptance tests to validate that systems do what the customers expect. Unit testing takes another view. Unit tests are written to ensure that code performs as the *programmer* expects.

Unit tests are generally focused at a lower level than other testing, establishing that underlying features work as expected. For example, an acceptance test might walk a user through an entire purchase. A unit test might verify that a `ShoppingCart` class correctly defends against adding an item with a negative quantity.

Unit testing is an example of *white box testing*, where knowledge of internal structures is used to identify the best ways to test the system. This is a complementary approach to *black box testing*, where the focus is not on implementation details but on overall functionality compared to specifications. You should leverage both approaches to effectively test your applications.

Benefits of unit testing

A common reaction to unit testing is to resist the approach because the tests seemingly make more work for a developer. However, unit testing offers many benefits that may not be obvious at first.

The act of writing tests often uncovers design or implementation problems. The unit tests serve as the first users of your system and will frequently identify design issues or functionality that is lacking. Once a unit test is written, it serves as a form of documentation for the use of the target system. Other developers can look to an assembly's unit tests to see example calls into various classes and members.

Perhaps one of the most important benefits is that a well-written test suite provides the original developer with the freedom to pass the system off to other developers for maintenance and further enhancement. Should those developers introduce a bug in the original functionality, there is a strong likelihood that those unit tests will detect that failure and help diagnose the issue. Meanwhile, the original developer can focus on current tasks.

It takes the typical developer time and practice to become comfortable with unit testing. Once a developer has been saved enough time by unit tests, he or she will latch on to them as an indispensable part of the development process.

Unit testing does require more explicit coding, but this cost will be recovered, and typically exceeded, when you spend much less time debugging your application. In addition, some of this cost is typically already hidden in the form of test console- or Windows-based applications. Unlike these informal testing applications, which are frequently discarded after initial verification, unit tests become a permanent part of the project, run each time a change is made to help ensure that the system still functions as expected. Tests are stored in source control very near to the code they verify and are maintained along with the code under test, making it easier to keep them synchronized.

Unit tests are an essential element of regression testing. Regression testing involves retesting a piece of software after new features have been added to make sure that errors or bugs are not introduced. Regression testing also provides an essential quality check when you introduce bug fixes in your product.

It is difficult to overstate the importance of comprehensive unit test suites. They enable a developer to hand off a system to other developers with confidence that any changes they make should not introduce undetected side effects. However, because unit testing only provides one view of a system's behavior, no amount of unit testing should ever replace integration, acceptance, and load testing.

Writing effective unit tests

Because unit tests are themselves code, you are generally unlimited in the approaches you can take when writing them. However, we recommend that you follow some general guidelines:

- ❑ Always separate your unit test assemblies from the code you are testing. This separation enables you to deploy your application code without unit tests, which serve no purpose in a production environment.
- ❑ Avoid altering the code you are testing solely to allow easier unit testing. A common mistake is to open accessibility to class members to allow unit tests direct access. This compromises design, reduces encapsulation, and broadens interaction surfaces. You will see later in this chapter that Team System offers features to help address this issue.
- ❑ Each test should verify a small slice of functionality. Do not write long sequential unit tests that verify a large number of items. While creating focused tests will result in more tests, the overall suite of tests will be easier to maintain. In addition, identifying the cause of a problem is much easier when you can quickly look at a small failed unit test, immediately understand what it was testing, and know where to search for the bug.
- ❑ All tests should be autonomous. Avoid creating tests that rely on other tests to be run beforehand. Tests should be executable in any combination and in any order. To verify that your tests are correct, try changing their execution order and running them in isolation.
- ❑ Test both expected behavior (normal workflows) and error conditions (exceptions and invalid operations). This often means that you will have multiple unit tests for the same method, but remember that developers will always find ways to call your objects that you did not intend. Expect the unexpected, code defensively, and test to ensure that your code reacts appropriately.

The final proof of your unit testing's effectiveness will be when they save you more time during development and maintenance than you spent creating them. In our experience, you will realize this savings many times over.

Test-driven development

Test-driven development (TDD) is the practice of writing unit tests before writing the code that will be tested. The logic behind writing tests against code that does not exist is difficult to grasp at first. Our experience has shown that TDD can be a challenge to introduce to developers and that the best way to learn it is by applying it on a small sample or noncritical project. Only after working with TDD will its advantages be fully understood and appreciated.

TDD encourages following a continuous cycle of development involving small and manageable steps. Figure 14-1 illustrates the steps involved in test-driven development.

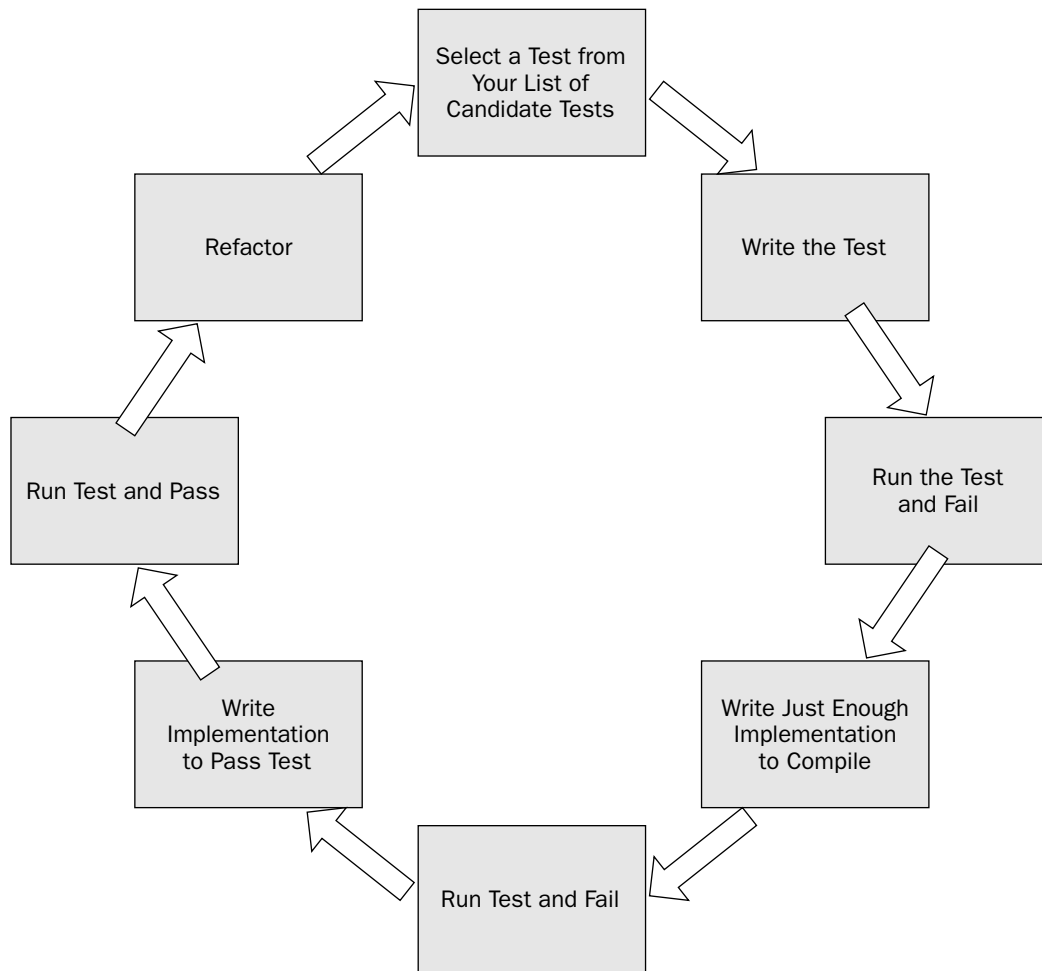


Figure 14-1

First, you generate an initial list of tests that apply to your task. Tests should verify all expected functionality, a variety of inputs, and error handling. Your initial list will rarely, if ever, include all of the tests you will need to write. As you work, you'll realize other inputs or scenarios that need testing. As this happens, simply add them to the list. Think of this list of tests as the requirements of your application. If your code needs to implement a certain behavior, make sure you have one or more unit tests to enforce it.

Once your initial list is ready, you now enter the cycle of TDD, shown in Figure 14-1. As shown at the top of the diagram, select one test from the list and begin writing the unit test. Of course, this cannot compile because the code you're testing does not exist. Write just enough application code to successfully compile and then run the test. As expected, it will fail because the implementation isn't complete. Add enough implementation logic to allow the test to pass. Resist the temptation to write all of the code you think will eventually be needed; instead, focus only on the current test. The code will be completed later as you implement all of the tests.

Once you have a passing test, you can refactor your code. As you learned in Chapter 11, refactoring is the process of making small changes to code that improve the design of the code but do not alter its functionality. Refactoring is critical when following TDD, as your code might not be as clean or organized as it might otherwise be when following a traditional design-first, top-down development model. Refactor to remove duplication and reduce complexity, and to improve maintainability and overall design.

Make sure your refactoring changes are done in small steps, building and running tests with each change. Once refactoring is complete, resume your cycle at the top by selecting the next test. When the list of tests is empty, you're ready to move on to your next task, confident that the functionality you've implemented has been tested and will be much easier to maintain.

This process is often abbreviated to what is known as "red, green, refactor." The "red" means you begin with a failing test. Once the implementation code is written, the test will pass, giving you the "green" result. Finally, the code is refactored to improve design.

You know from Chapter 11 that Visual Studio helps you perform refactoring in efficient and less error-prone ways. Later in this chapter, you'll see how Team System can facilitate a TDD approach through its integrated unit testing features and by offering code generation from within unit tests.

Third-party tools

Unit testing is not a new concept. Before Team System introduced integrated unit testing, developers needed to rely on third-party frameworks. The de facto standard for .NET unit testing has been an open-source package called *NUnit*. NUnit has its original roots as a .NET port of the Java-based JUnit unit testing framework. JUnit is itself a member of the extended xUnit family.

There are many similarities between NUnit and the unit testing framework in Team System. The structure and syntax of tests and the execution architecture are conveniently similar. If you have existing suites of NUnit-based tests, it is generally easy to convert them for use with Team System.

James Newkirk, who led development of NUnit 2.0 and is now employed at Microsoft, has written a tool to convert your NUnit tests to Team System's syntax. At the time of writing, this tool, called the "NUnit Converter," is available from GotDotNet at <http://workspaces.gotdotnet.com/nunitaddons>.

Team System's implementation of unit testing is not merely a port of NUnit. Microsoft has added a number of features that are unavailable with the version of NUnit available at the time of this writing. Among these are IDE integration, code generation, new attributes, enhancements to the `Assert` class, and built-in support for testing nonpublic members. We describe all of these in detail later in this chapter.

Team System Unit Testing

Unit testing is a feature available in the Developer and Tester editions of Visual Studio Team System. In this section, we describe how to create, execute, and manage unit tests.

Unit tests are themselves normal code, identified as unit tests through the use of attributes. Like NUnit 2.0 and later, Team System uses .NET reflection to inspect assemblies to find unit tests.

Chapter 14

Reflection is a mechanism by which details about .NET objects can be discovered at execution time. The `System.Reflection` assembly contains members that help you identify classes, properties, and methods of any .NET assembly. Reflection even enables you to call methods and access properties of classes. This includes access to private members, a practice that can be useful in unit testing, as you will see later in this chapter.

You will also use attributes to identify other structures used in your tests and to indicate desired behaviors.

Creating your first unit test

In this section, we'll take a slower approach to creating a unit test than you will in your normal work. This will give you a chance to examine details you could miss using only the built-in features that make unit testing easier. Later in this chapter, you'll look at the faster approaches.

In order to have something to test, create a new C# Class Library project named `ExtendedMath`. Rename the default `Class1.cs` to `Functions.cs`. We'll add code to compute the Fibonacci for a given number. The Fibonacci Sequence, as you may recall, is a series of numbers where each term is the sum of the prior two terms. The first six terms, starting with an input factor of 1, are {1, 1, 2, 3, 5, 8}.

Open `Functions.cs` and insert the following code:

```
using System;

namespace ExtendedMath
{
    public static class Functions
    {
        public static int Fibonacci(int factor)
        {
            if (factor < 2)
                return (factor);
            int x = Fibonacci(--factor);
            int y = Fibonacci(--factor);

            return x + y;
        }
    }
}
```

You are now ready to create unit tests to verify the `Fibonacci` implementation. Unit tests are recognized as tests only if they are contained in separate projects called *test projects*. Test projects can contain any of the test types supported in Team System. Add a test project named `ExtendedMathTesting` to your solution by adding a new project and selecting the Test Project template. If the test project includes any sample tests for you, such as `UnitTest1.cs` or `ManualTest1.mht`, you can safely delete them.

For full details on creating and customizing test projects, see Chapter 13.

Because you will be calling objects in your `ExtendedMath` project, make a reference to that class library project from the test project. You may notice that a reference to the `Microsoft.VisualStudio.TestTools.UnitTesting.dll` assembly has already been made for you. This assembly contains many helpful classes for creating units tests. We'll use many of these throughout this chapter.

As you'll see later in this chapter, Team System supports the generation of basic unit test outlines, but in this section, we'll create them manually to make it easier to understand the core concepts.

Once you have created a new test project, add a new class file (not a unit test; we'll cover that file type later) called `FunctionsTest.cs`. You will use this class to contain the unit tests for the `Functions` class. You'll be using unit testing objects from the `ExtendedMath` project and the `UnitTestFixture` assembly mentioned earlier, so add `using` statements at the top so that the class members do not need to be fully qualified:

```
using ExtendedMath;
using Microsoft.VisualStudio.TestTools.UnitTesting;
```

Identifying unit test classes

To enable Team System to identify a class as potentially containing unit tests, you must assign the `TestClass` attribute. If you forget to add the `TestClass` attribute, the unit tests methods in your class will not be recognized.

To indicate that the `FunctionsTest` class will contain unit tests, add the `TestClass` attribute to its declaration:

```
namespace ExtendedMath
{
    [TestClass]
    public class FunctionsTest
    {
    }
}
```

Unit tests are required to be hosted within public classes, so don't forget to include the `public` descriptor for the class. Note also that parentheses after an attribute are optional if you are not passing parameters to the attribute. For example, `[TestClass()]` and `[TestClass]` are equivalent.

Identifying unit tests

Having identified the class as a container of unit tests, you're ready to add your first unit test. A unit test method must be public, nonstatic, accept no parameters, and have no return value. To differentiate unit test methods from ordinary methods, they must be decorated with the `TestMethod` attribute.

Add the following code inside the `FunctionsTest` class:

```
[TestMethod]
public void FibonacciTest()
{
}
```

Chapter 14

Unit test success and failure

You have the shell of a unit test, but how do you test? A unit test indicates failure to Team System by throwing an exception. Any test that does not throw an exception is considered to have passed, except in the case of `ExpectedException` attribute, which we describe later.

The unit testing framework defines the `Assert` object. This object exposes many members, which are central to creating unit tests. You'll learn more about `Assert` later in the chapter.

Add the following code to the `FibonacciTest`:

```
[TestMethod]
public void FibonacciTest()
{
    const int FACTOR = 8;
    const int EXPECTED = 21;

    int actual = ExtendedMath.Functions.Fibonacci(FACTOR);

    Assert.AreEqual(EXPECTED, actual);
}
```

This uses the `Assert.AreEqual` method to compare two values, the value you expect and the value generated by calling the `Fibonacci` method. If they do not match, an exception will be thrown, causing the test to fail.

When you run tests, you will see the Test Results window. Success is indicated with a green checkmark, failure with a red X. A special result, inconclusive (described later in this chapter), is represented by a question mark.

To see a failing test, change the `EXPECTED` constant from 21 to 22 and rerun the test. The Test Results window will show the test as failed. The Error Message column provides details about the failure reason. In this case, the Error Message would show the following:

```
Assert.AreEqual failed. Expected:<22>, Actual:<21>
```

This indicates that either the expected value is wrong or the implementation of the `Fibonacci` algorithm is wrong. Fortunately, because unit tests verify a small amount of code, the job of finding the source of bugs is made easier.

Managing and running unit tests

Once you have created a unit test and rebuilt your project, Visual Studio will automatically inspect your projects for unit tests. You can use the Test Manager and Test View windows to work with your tests.

For full details on test management and related windows, see Chapter 13.

The easiest way to open these windows is by enabling the Test Tools toolbar and pressing either the Test View or Test Manager buttons. They are also available by selecting `Test ⇄ Windows`.

Test View

Test View provides a compact view of your tests. It enables you to quickly select and run your tests. You can group tests by name, project, type, class name, and other criteria. Figure 14-2 shows the Test View window.

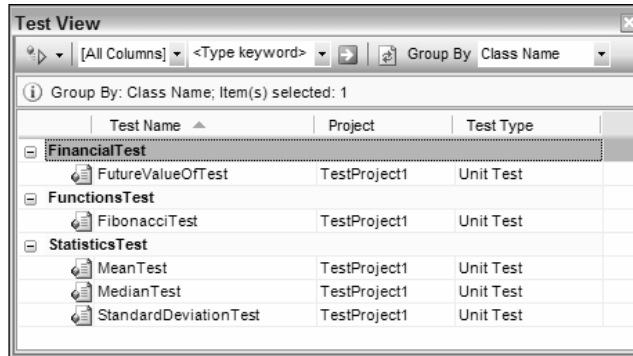


Figure 14-2

Double-click on any test to navigate to that test's code. To run one or more tests, select them and press the Run Selection button.

Test Manager

The Test Manager offers all of the features of the Test View window, but provides more options for organization and display of your tests. Figure 14-3 shows the Test Manager.

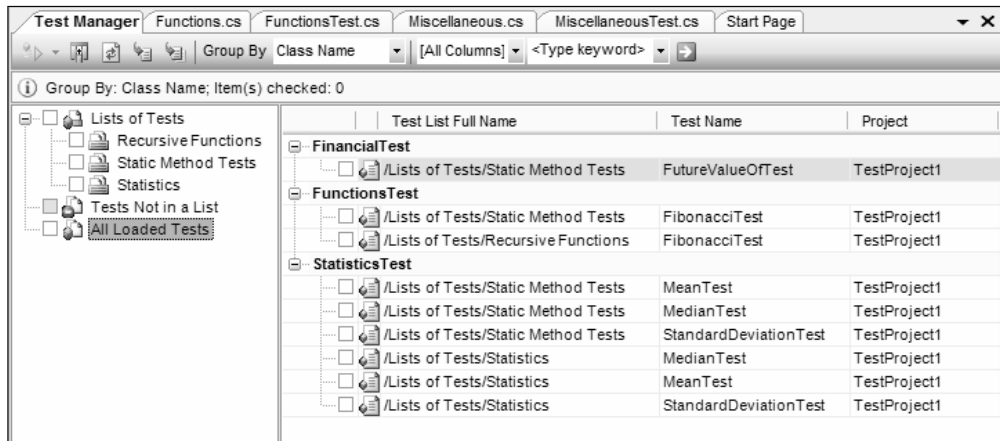


Figure 14-3

By default, all tests are listed, but you can organize tests into lists. In Test Manager, right-click on the Lists of Tests node on the tree on the left-hand side. Select New Test List and the Create New Test List dialog will appear as shown in Figure 14-4. You can also create a new test list by choosing Test ⇨ Create New Test List.

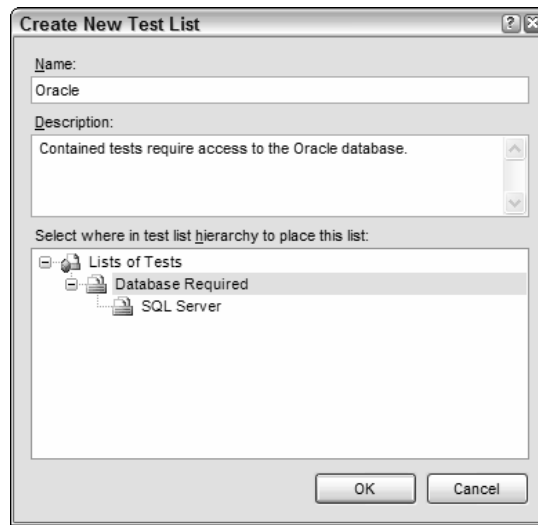


Figure 14-4

Give the list a name and optionally a description. You can also place this list within another list. For example, as shown in Figure 14-4, you might have a top-level Database Required list with a sublist of SQL Server, to which a second sublist for Oracle is being added.

Once you have created a list, you can drag tests from the Test Manager list onto it. The list enables you to easily select all related tests for execution by checking the box next to the list name.

You can also group by test properties. First, click the All Loaded Tests node to ensure that all tests are listed. Then, in the Group By drop-down list, scroll to and select Class Name. You will then see your tests in collapsible groups by class name.

If you have many tests, you will find the filtering option useful. With filters, you can limit the list of tests to only those that match text you enter. Next to the Group By list is a Filter Column drop-down and a text box. Enter the text you wish to match in the text box; you can optionally select a column you would like to match against. For example, if you wish to show only tests from a certain class, enter that class name then select Class Name in the Filter Column list. When you press the Apply Filter button, only tests from that class will be displayed. If the Filter Column is set to [All Columns], tests will display if any of their fields contain the text you enter.

Test run configuration

Whenever you run a set of tests, a group of settings apply to that run. Those settings, called the *test run configuration*, are stored in an XML file with a `.testrunconfig` extension. A test run configuration is created automatically for every new test project, named `localtestrun.testrunconfig`.

The settings include items such as the naming structure of your results files, configuration of remote test execution, enabling of code coverage, and specifying additional files to deploy during your tests.

To edit a test run configuration, choose Test ⇨ Edit Test Run Configurations, and then choose the configuration you wish to modify. You can also double-click the configuration file in Solution Explorer. Figure 14-5 shows the Test Run Configuration interface.

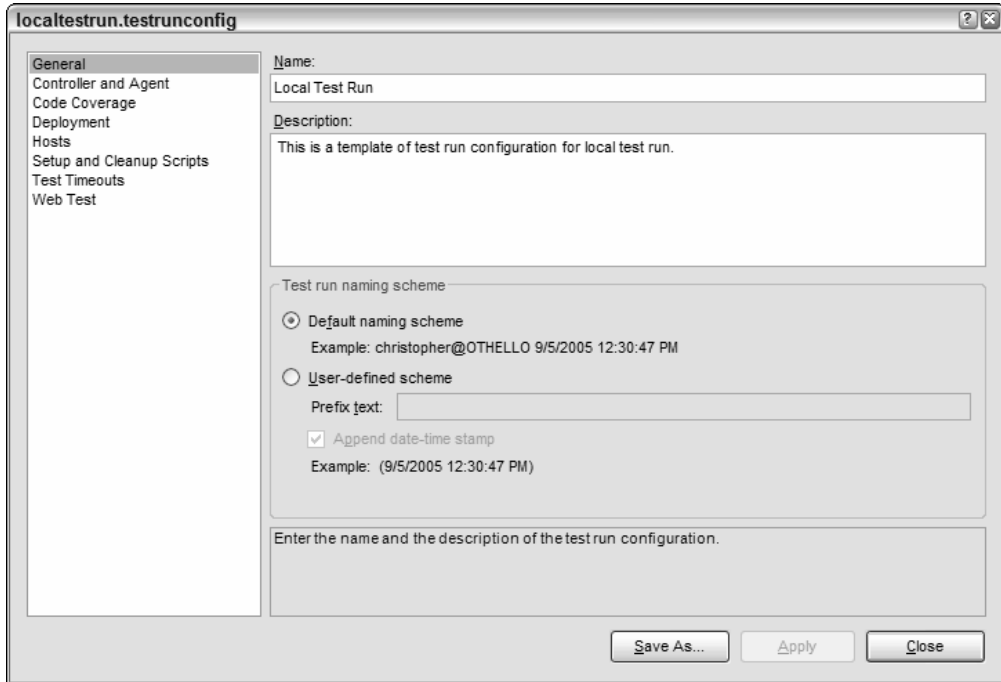


Figure 14-5

You may have more than one test run configuration, perhaps to support different execution environments or code coverage settings, but you must select a single configuration as “active” when you run your tests. To set another configuration as active, choose Test ⇨ Select Active Test Run Configuration, and then choose the correct configuration.

The default configuration settings will generally be fine for your initial unit testing. As you begin to require additional unit testing features, you may need to adjust these settings. For example, later in this chapter, we describe how to monitor code coverage with Team System. In that section, you will learn how to use the test run configuration to enable that feature.

Test results

Once you have selected and run one or more unit tests using the Test View or Test Manager windows, you will see the Test Results window. This window displays the status of your test execution, as shown in Figure 14-6.

You can see that one of the tests has failed, two have passed, and two are inconclusive. The error message for any nonpassing test is displayed. You can double-click on any test result to see details for that test. You can also right-click on a test result and choose Open Test to navigate to the unit test code.

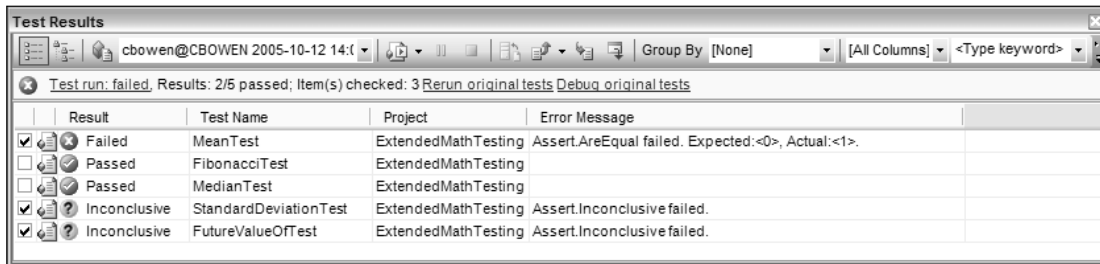


Figure 14-6

Notice that the nonpassing tests are checked. This convenient feature enables you to make some changes and then quickly rerun just those tests that have not passed.

Debugging unit tests

Because unit tests are simply methods with special attributes applied to them, they can be debugged just like other code.

Breakpoints can be set anywhere in your code, not just in your unit tests. For example, the `FibonacciTest` calls into the `ExtendedMath.Fibonacci` method. You could set a breakpoint in either method and have execution paused when that line of code is reached.

However, setting program execution will not pause at your breakpoints unless you run your unit test in debugging mode. The Test View, Test Results, and Test Manager windows all feature a drop-down arrow next to the Run button. For example, in Test Manager, click on the arrow next to the Run Checked Tests button. You will see a new option, Debug Checked Tests. If you choose this, the selected unit tests will be run in debug mode, pausing execution at any enabled breakpoints, and giving you a chance to evaluate and debug your unit test or implementation code as necessary.

If you have enabled code coverage for your application, you will see a message indicating that you cannot debug while code coverage is enabled. Click OK and you will continue debugging as normal, but code coverage results will not be available. We describe code coverage later in this chapter.

Keep in mind that the Run/Debug buttons are “sticky.” The last selected mode will continue to be used when the button is pressed until another mode is chosen from the drop-down list.

Programming with the Unit Test Framework

In this section, we describe in detail the attributes and methods available to you for creating unit tests. All of the classes and attributes mentioned in this section can be found in the `Microsoft.VisualStudio.TestTools.UnitTesting` namespace.

Initialization and cleanup of unit tests

Often, you’ll need to configure a resource that is shared among your tests. Examples might be a database connection, a log file, or a shared object in a known default state. You might also need ways to clean up from the actions of your tests, such as closing a shared stream or rolling back a transaction.

The unit test framework offers attributes to identify such methods. They are grouped into three levels: Test, Class, and Assembly. The levels determine the scope and timing of execution for the methods they decorate. The following table describes these attributes.

Attributes	Frequency and Scope
TestInitialize, TestCleanup	Executed before (Initialize) or after (Cleanup) any of the class's unit tests are run
ClassInitialize, ClassCleanup	Executed a single time before or after any of the tests in the current class are run
AssemblyInitialize, AssemblyCleanup	Executed a single time before or after any number of tests in any of the assembly's classes are run

Having methods with these attributes is optional, but do not define more than one of each attribute in the same context.

Do not use class- or assembly-level initialize and cleanup attributes with ASP.NET unit tests. When run under ASP.NET, these methods cannot be guaranteed to run only once. Because these are static methods, this may lead to false testing results.

TestInitialize and TestCleanup attributes

Use the `TestInitialize` attribute to create a method that will be executed one time before each unit test run in the current class. Similarly, `TestCleanup` marks a method that will always run immediately after each test. Like unit tests, methods with these attributes must be public, nonstatic, accept no parameters, and have no return values.

Here is an example test for a simplistic shopping cart class. It contains two tests and defines the `TestInitialize` and `TestCleanup` methods:

```
using Microsoft.VisualStudio.TestTools.UnitTesting;
```

```
[TestClass]
public class ShoppingCartTest
{
    private ShoppingCart cart;

    [TestInitialize]
    public void TestInitialize()
    {
        cart = new SomeClass();
        cart.Add(new Item("Test"));
    }
}
```

```
[TestCleanup]
public void TestCleanup()
{
    // Not required - here for illustration
    cart.Dispose();
}
```

Chapter 14

```
    }

    [TestMethod]
    public void TestCountAfterAdd()
    {
        int expected = cart.Count + 1;
        cart.Add(new Item("New Item"));
        Assert.AreEqual(expected, cart.Count);
    }

    [TestMethod]
    public void TestCountAfterRemove()
    {
        int expected = cart.Count - 1;
        cart.Remove(0);
        Assert.AreEqual(expected, cart.Count);
    }
}
```

When you run both tests, `TestInitialize` and `TestCleanup` are both executed twice. `TestInitialize` is run immediately before each unit test and `TestCleanup` immediately after.

ClassInitialize and ClassCleanup attributes

The `ClassInitialize` and `ClassCleanup` attributes are used very similarly to `TestInitialize` and `TestCleanup`. The difference is that these methods are guaranteed to run once and only once no matter how many unit tests are executed from the current class. Unlike `TestInitialize` and `TestCleanup`, these methods are marked static and accept a `TestContext` instance as a parameter.

The importance of the `TestContext` instance is described later in this chapter.

The following code demonstrates how you might manage a shared logging target using class-level initialization and cleanup with a logging file:

```
private System.IO.File logFile;

[ClassInitialize]
public static void ClassInitialize(TestContext context)
{
    // Code to open the logFile object
}

[ClassCleanup]
public static void ClassCleanup(TestContext context)
{
    // Code to close the logFile object
}
```

You could now reference the `logFile` object from any of your unit tests in this class, knowing that it will automatically be opened before any unit test is executed and closed after the final test in the class has completed.

This approach to logging is simply for illustration. You'll see later how the `TestContext` object passed into these methods enables you to more effectively log details from your unit tests.

The following code shows the flow of execution should you again run both tests:

```
ClassInitialize
  TestInitialize
    TestCountAfterAdd
  TestCleanup
  TestInitialize
    TestCountAfterRemove
  TestCleanup
ClassCleanup
```

AssemblyInitialize and AssemblyCleanup attributes

Where you might use `ClassInitialize` and `ClassCleanup` to control operations at a class level, use the `AssemblyInitialize` and `AssemblyCleanup` attributes for an entire assembly. For example, a method decorated with `AssemblyInitialize` will be executed once before any test in that current assembly, not just those in the current class. As with the class-level initialize and cleanup methods, these must be static and accept a `TestContext` parameter:

```
[AssemblyInitialize]
public static void AssemblyInitialize(TestContext context)
{
    // Assembly-wide initialization code
}
```

```
[AssemblyCleanup]
public static void AssemblyCleanup(TestContext context)
{
    // Assembly-wide cleanup code
}
```

Consider using `AssemblyInitialize` and `AssemblyCleanup` in cases where you have common operations spanning multiple classes. Instead of having many per-class initialize and cleanup methods, you can refactor these to single assembly-level methods.

Using the Assert methods

The most common way to determine success in unit tests is to compare an expected result against an actual result. The `Assert` class features many methods that enable you to make these comparisons quickly.

Assert.AreEqual and Assert.AreNotEqual

Of the various `Assert` methods, you will likely find the most use for `AreEqual` and `AreNotEqual`. As their names imply, you are comparing an expected value to a supplied value. If the operands are not value-equivalent (or are equivalent for `AreNotEqual`), then the current test will fail.

A third, optional argument can be supplied: a string that will be displayed along with your unit test results, which you can use to describe the failure. Additionally, you can supply parameters to be replaced in the string, just as the `String.Format` method supports:

Chapter 14

```
[TestMethod]
public void IsPrimeTest()
{
    const int FACTOR = 5;
    const bool EXPECTED = true;

    bool actual = CustomMath.IsPrime(FACTOR);

    Assert.AreEqual(EXPECTED, actual, "The number {0} should have been computed as
        prime, but was not.", FACTOR);
}
```

`Assert.AreEqual` and `Assert.AreNotEqual` have many parameter overloads, accepting types such as `string`, `double`, `int`, `float`, `object`, and generic types. Take the time to review the overloads in the [Object Browser](#).

When using these methods with two string arguments, one of the overrides allows you to optionally supply a third argument. This is a Boolean, called `ignoreCase`, that indicates whether the comparison should be case-insensitive. The default comparison is case-sensitive.

Working with floating-point numbers involves a degree of imprecision. You can supply an argument that defines a delta by which two numbers can differ yet still pass a test—for example, if you’re computing square roots and decide that a “drift” of plus or minus 0.0001 is acceptable:

```
[TestMethod]
public void SquareRootTest()
{
    const double EXPECTED = 3.1622;
    const double DELTA = 0.0001;
    double actual = CustomMath.SquareRoot(10);

    Assert.AreEqual(EXPECTED, actual, DELTA, "Root not within acceptable range");
}
```

Assert.AreSame and Assert.AreNotSame

`AreSame` and `AreNotSame` function in much the same manner as `AreEqual` and `AreNotEqual`. The important difference is that these methods compare the *references* of the supplied arguments. For example, if two arguments point to the same object instance, then `AreSame` will pass. Even when the arguments are exactly equivalent in terms of their state, `AreSame` will fail if they are not in fact the same object. This is the same concept that differentiates `object.Equals` from `object.ReferenceEquals`.

A common use for these methods is to ensure that properties return expected instances or that collections handle references correctly. In the following example, we add an item to a collection and ensure that what we get back from the collection’s indexer is a reference to the same item instance:

```
[TestMethod]
public void CollectionTest()
{
    CustomCollection cc = new CustomCollection();
    Item original = new Item("Expected");
    cc.Add(original);
```



```
Item actual = cc[0];

Assert.AreSame(original, actual);
}
```

Assert.IsTrue and Assert.IsFalse

As you can probably guess, `IsTrue` and `IsFalse` are used simply to ensure that the supplied expression is true or false as expected. Returning to the `IsPrimeNumberTest` example, we can restate it as follows:

```
[TestMethod]
public void IsPrimeTest()
{
    const int FACTOR = 5;

    Assert.IsTrue(CustomMath.IsPrime(FACTOR), "The number {0} should have been
        computed as prime, but was not.", FACTOR);
}
```

Assert.IsNull and Assert.IsNotNull

Similar to `IsTrue` and `IsFalse`, these methods verify that a given object type is either null or not null. Revising the collection example, this ensures that the item returned by the indexer is not null:

```
[TestMethod]
public void CollectionTest()
{
    CustomCollection cc = new CustomCollection();
    cc.Add(new Item("Added"));
    Item item = cc[0];

    Assert.IsNotNull(item);
}
```

Assert.IsInstanceOfType and Assert.IsNotInstanceOfType

`IsInstanceOfType` simply ensures that a given object is an instance of an expected type. For example, suppose you have a collection that accepts entries of any type. You'd like to ensure that an entry you're retrieving is of the expected type:

```
[TestMethod]
public void CollectionTest()
{
    UntypedCollection untyped = new UntypedCollection();
    untyped.Add(new Item("Added"));
    untyped.Add(new Person("Rachel"));
    untyped.Add(new Item("Another"));

    object entry = untyped[1];

    Assert.IsInstanceOfType(entry, typeof(Person));
}
```

As you can no doubt guess, `IsNotInstanceOfType` will test to ensure that an object is not the specified type.

Assert.Fail and Assert.Inconclusive

Use `Assert.Fail` to immediately fail a test. For example, you may have a conditional case that should never occur. If it does, call `Assert.Fail` and an `AssertFailedException` will be thrown, causing the test to abort with failure. You may find `Assert.Fail` useful when defining your own custom `Assert` methods.

`Assert.Inconclusive` enables you to indicate that the test result cannot be verified as a pass or fail. This is typically a temporary measure until a unit test (or the related implementation) has been completed. As described in the section “Code Generation” later in this chapter, `Assert.Inconclusive` is used to indicate that more work is needed to be done to complete a unit test.

There is no `Assert.Succeed` because success is indicated by completion of a unit test method without a thrown exception. Use a return statement if you wish to cause this result from some point in your test.

`Assert.Fail` and `Assert.Inconclusive` both support a string argument and optional arguments, which will be inserted into the string in the same manner as `String.Format`. Use this string to supply a detailed message back to the Test Results window, describing the reasons for the nonpassing result.

Using the CollectionAssert class

The `Microsoft.VisualStudio.TestTools.UnitTesting` namespace includes a class, `CollectionAssert`, containing useful methods for testing the contents and behavior of collection types.

The following table describes the methods supported by `CollectionAssert`.

Method	Description
<code>AllItemsAreInstancesOfType</code>	Ensures that all elements are of an expected type
<code>AllItemsAreNotNull</code>	Ensures that no items in the collection are <code>null</code>
<code>AllItemsAreUnique</code>	Searches a collection, failing if a duplicate member is found
<code>AreEqual</code>	Ensures that two collections have reference-equivalent members
<code>AreNotEqual</code>	Ensures that two collections do not have reference-equivalent members
<code>AreEquivalent</code>	Ensures that two collections have value-equivalent members
<code>AreNotEquivalent</code>	Ensures that two collections do not have value-equivalent members
<code>Contains</code>	Searches a collection, failing if the given object is not found
<code>DoesNotContain</code>	Searches a collection, failing if a given object is found
<code>IsNotSubsetOf</code>	Ensures that the first collection has members not found in the second
<code>IsSubsetOf</code>	Ensures that all elements in the first collection are found in the second

The following example uses some of these methods to verify various behaviors of a collection type, `CustomCollection`. When this example is run, none of the assertions fail and the test results in success. Note that proper unit testing would spread these checks across multiple smaller tests.

```
[TestMethod]
public void CollectionTests()
{
    CustomCollection list1 = new CustomCollection();
    list1.Add("alpha");
    list1.Add("beta");
    list1.Add("delta");
    list1.Add("delta");

    CollectionAssert.AllItemsAreInstancesOfType(list1, typeof(string));
    CollectionAssert.AllItemsAreNotNull(list1);

    CustomCollection list2 = (CustomCollection)list1.Clone();

    CollectionAssert.AreEqual(list1, list2);
    CollectionAssert.AreEquivalent(list1, list2);

    CustomCollection list3 = new CustomCollection();
    list3.Add("beta");
    list3.Add("delta");

    CollectionAssert.AreNotEquivalent(list3, list1);
    CollectionAssert.IsSubsetOf(list3, list1);
    CollectionAssert.DoesNotContain(list3, "alpha");
    CollectionAssert.AllItemsAreUnique(list3);
}
```

The final assertion, `AllItemsAreUnique(list3)`, would have failed if tested against `list1` because that collection has two entries of the string "delta".

Using the *StringAssert* class

Similar to `CollectionAssert`, the `StringAssert` class contains methods that enable you to easily make assertions based on common text operations. The following table describes the methods supported by `StringAssert`.

Method	Description
Contains	Searches a string for a substring and fails if not found
DoesNotMatch	Applies a regular expression to a string and fails if any matches are found
EndsWith	Fails if the string does not end with a given substring
Matches	Applies a regular expression to a string and fails if no matches are found
StartsWith	Fails if the string does not begin with a given substring

Chapter 14

Here are some simple examples of these methods. Each of these assertions will pass:

```
[TestMethod]
public void TextTests()
{
    StringAssert.Contains("This is the searched text", "searched");

    StringAssert.EndsWith("String which ends with searched", "ends with searched");

    StringAssert.Matches("Search this string for whitespace",
        new System.Text.RegularExpressions.Regex(@"\s+"));

    StringAssert.DoesNotMatch("Doesnotcontainwhitespace",
        new System.Text.RegularExpressions.Regex(@"\s+"));

    StringAssert.StartsWith("Starts with correct text", "Starts with");
}
```

`Matches` and `DoesNotMatch` accept a string and an instance of `System.Text.RegularExpressions.Regex`. In the preceding example, a simple regular expression that looks for at least one whitespace character was used. `Matches` finds whitespace and the `DoesNotMatch` does not find whitespace, so both pass.

Expecting exceptions

Normally, a unit test that throws an exception is considered to have failed. However, you'll often wish to verify that a class behaves correctly by throwing an exception. For example, you might provide invalid arguments to a method to verify that it properly throws an exception.

The `ExpectedException` attribute indicates that a test will succeed only if the indicated exception is thrown. Not throwing an exception or throwing an exception of a different type will result in test failure.

The following unit test expects that an `ObjectDisposedException` will be thrown:

```
[TestMethod]
[ExpectedException(typeof(ObjectDisposedException))]
public void ReadAfterDispose()
{
    CustomFileReader cfr = new CustomFileReader("target.txt");
    cfr.Dispose();
    string contents = cfr.Read(); // Should throw ObjectDisposedException
}
```

The `ExpectedException` attribute supports a second, optional string argument. The `Message` property of the thrown exception must match this string or the test will fail. This enables you to differentiate between two different instances of the same exception type.

For example, suppose you are calling a method that throws a `FileNotFoundException` for several different files. To ensure that it cannot find one specific file in your testing scenario, supply the message you expect as the second argument to `ExpectedException`. If the exception thrown is not `FileNotFoundException` and its `Message` property does not match that text, the test will fail.

Defining custom unit test properties

You may define custom properties for your unit tests. For example, you may wish to specify the author of each test and be able to view that property from the Test Manager.

Use the `TestProperty` attribute to decorate a unit test, supplying the name of the property and a value:

```
[TestMethod]
[TestProperty("Author", "Deborah")]
public void ExampleTest()
{
    // Test logic
}
```

Now, when you view the properties of that test, you will see a new entry, `Author`, with the value `Deborah`. If you change that value from the Properties window, the attribute in your code will automatically be updated.

TestContext class

Unit tests normally have a reference to a `TestContext` instance. This object provides run-time features that might be useful to tests, such as details of the test itself, the various directories in use, and several methods to supplement the details stored with the test's results. `TestContext` is also very important for data-driven and ASP.NET unit tests, as you will see later.

Several methods are especially useful to all unit tests. The first, `WriteLine`, enables you to insert text into the results of your unit test. This can be useful for supplying additional information about the test, such as parameters, environment details, and other debugging data that would normally be excluded from test results. By default, information from the test run is stored in a *test results file*, an XML file with a `.trx` extension. These files can be found in the `TestResults` subdirectory of your project. The default name for the files is based on the user, machine, and date of the test run, but this can be modified via the test run configuration settings. See Chapter 13 for more information on test results files.

Here is a simple example of a unit test that accesses the `TestContext` to send a string containing the test's name to the results:

```
[TestClass]
public class TestClass
{
    private TestContext testContextInstance;

    public TestContext TestContext
    {
        get { return testContextInstance; }
        set { testContextInstance = value; }
    }

    [TestMethod]
    public void TestMethod1()
    {
        TestContext.WriteLine("This is test {0}", TestContext.TestName);
    }
}
```

Chapter 14

The `AddResultFile` method enables you to add a file, at runtime, to the results of the test run. The file you specify will be copied to the results directory alongside other results content. For example, this may be useful if your unit test is validating an object that creates or alters a file and you would like that file to be included with the results for analysis.

Finally, the `BeginTimer` and `EndTimer` methods enable you to create one or more named timers within your unit tests. The results of these timers are stored in the test run's results.

Creating data-driven unit tests

An excellent way to verify the correct behavior of code is to execute it using realistic data. Team System provides features to automatically bind data from a data source as input to unit tests. The unit test is run once for each data row.

A unit test is made data-driven by assigning attributes for connecting to and reading from a data source. The easiest way to do this is to modify an existing unit test's properties in the Test Manager window. Begin with a normal unit test outline, and then open Test Manager. Select the unit test and view its properties.

First, establish the connection to the data source by setting the Data Connection String property. You may either enter it manually or click the button labeled with ellipses ("...") to use dialogs to create the connection string. Once the connection is specified, select the table you wish to use in the Data Table Name property.

As mentioned before, the unit test will be called once per row in the data source. You can set how rows are fed into the unit test via the Data Access Method property. Sequential will feed the rows in exactly the order returned from the data source, whereas Random will select random rows. Rows are provided to the unit test until all rows have been used once.

Setting these properties will automatically decorate the selected unit test with the appropriate attributes. You do not need to use the Properties window to create data-driven tests. For example, you may wish to copy attributes from an existing data-driven test to quickly create another.

To access the data from within the unit test, use the `DataRow` property of the `TestContext`. For example, if you bound to a table with customer data, you could read the current customer's ID from the `CustomerID` column with the following:

```
long customerID = TestContext.DataRow["CustomerID"];
```

Besides a column name, the `DataRow` property also accepts a column offset. If `CustomerID` were the first column in the table, you could supply a zero as the argument with the same result.

Because the unit test is run once per row in the data source, you want to be able to tell which rows caused certain results. Fortunately, this detail is already tracked for you. In the Test Results window, right-click on the test result and choose View Test Results Details. You will see a list of pass/fail results for each record in the database, enabling you to easily see which rows caused your test to fail.

Team System makes it very easy to write comprehensive tests against actual data. However, keep in mind that it is not enough to test only valid data, such as your real customer data. You will also want to have unit tests that verify your code's behavior when invalid data, perhaps in this case a negative number, is supplied.

Accessing Nonpublic Members from Tests

What if you want to test a class member that is not public? For example, suppose you're writing a private function that is never publicly accessed and is only used internally by other members. You'd like to test this method, but your test code does not have access to private members (or to internal members if the code is in a separate assembly).

There are three main approaches for addressing this issue:

- ☐ Make the private members you need to test public.
- ☐ Ensure that the private members are reachable through a public member and test via those public members.
- ☐ Use .NET reflection in the tests to load and directly invoke the nonpublic members.

With Team System, this final approach is abstracted for you. The following two sections describe how Team System helps to automate this previously manual task.

Testing private members is a controversial subject. Some people prefer to test only via public members to allow for easier refactoring. Others argue that an API should never be modified just for the sake of easier testing. If you agree with the former opinion, you can safely skip the remainder of this section.

Using PrivateObject to access nonpublic instance members

Suppose you'd like to test the private field and method of the following class:

```
public class Example
{
    public Example() {}

    private string password = "letmein";

    private bool VerifyPassword(string password)
    {
        return (String.Compare(this.password, password, false) == 0);
    }
}
```

Because the field and method are marked `private`, a unit test will not have the ability to directly access them. How can you ensure that `VerifyPassword` is working correctly?

Team System introduces the `PrivateObject` class, which is a wrapper around reflection code that enables you to access nonpublic members in a fairly straightforward manner.

The following table summarizes the methods supported by `PrivateObject`.

Chapter 14

Method	Description
GetArrayElement	Returns the selected item from a private array member. Supports multidimensional arrays with additional arguments.
GetField	Returns the value of the target field
GetFieldOrProperty	Returns the value of the target field or property
GetProperty	Returns the value of the target property
Invoke	Invokes the target method, optionally passing parameters
SetArrayElement	Assigns the given value to the indicated element of a private array Supports multidimensional arrays with additional arguments
SetField	Assigns the supplied object to the target field
SetFieldOrProperty	Assigns the supplied object to the target field or property
SetProperty	Assigns the supplied object to the target property

To use it, you first create an instance of the `PrivateObject` class, passing a `Type` object for the class you wish to work with:

```
using System;
using Microsoft.VisualStudio.TestTools.UnitTesting;

namespace Explorations
{
    [TestClass]
    public class ExampleTest
    {
        private PrivateObject privateObject;
        const string PASSWORD = "letmein";

        [TestInitialize]
        public void TestInitialize()
        {
            privateObject = new PrivateObject(typeof(Example));
        }
    }
}
```

Now you can create your tests. Use the `GetField` method of the `PrivateObject` instance to access non-public fields, supplying the name of the desired variable. Similarly, use `Invoke` to call methods, supplying the method name as the first argument, followed by any parameters to that method:

```
using System;
using Microsoft.VisualStudio.TestTools.UnitTesting;

namespace Explorations
{
    [TestClass]
    public class ExampleTest
    {

```



```
private PrivateObject privateObject;
const string PASSWORD = "letmein";

[TestInitialize]
public void TestInitialize()
{
    privateObject = new PrivateObject(typeof(Example));
}

[TestMethod]
public void ComparePrivatePassword()
{
    string password = (string)privateObject.GetField("password");
    Assert.AreEqual(PASSWORD, password);
}

[TestMethod]
public void TestPrivateVerifyPassword()
{
    bool accepted = (bool)privateObject.Invoke("VerifyPassword", PASSWORD);
    Assert.IsTrue(accepted);
}
}
```

Because `PrivateObject` uses reflection, you need to cast the results of these calls from the generic `Object` type back to the correct underlying type.

Using `PrivateType` to access nonpublic static members

`PrivateObject` is used to access instance-based members of a class. If you need to access static non-public members, you use the `PrivateType` class, which has a very similar interface and is a wrapper of reflection code.

The following table summarizes the methods exposed by `PrivateType`.

Method	Description
<code>GetStaticArrayElement</code>	Returns the selected item from a private static array member. Supports multidimensional arrays with additional arguments.
<code>GetStaticField</code>	Returns the value of the target static field.
<code>GetStaticProperty</code>	Returns the value of the target static property.
<code>InvokeStatic</code>	Invokes the target static method, optionally passing parameters.
<code>SetStaticArrayElement</code>	Assigns the given value to the indicated element of a private static array. Supports multidimensional arrays with additional arguments.
<code>SetStaticField</code>	Assigns the supplied object to the target static field.
<code>SetStaticProperty</code>	Assigns the supplied object to the target static property.

Chapter 14

The usage is very similar to the `PrivateObject`. Create an instance of the `PrivateType`, indicating which type you wish to work with, and then use the methods to access members as with `PrivateObject`. Suppose you added a private static count of password failures with a wrapping private property called `FailureCount`. The following code could read and test that property:

```
[TestMethod]
public void TestPrivateStaticFailureCount()
{
    PrivateType example = new PrivateType(typeof(Example));
    int failureCount = (int)example.GetStaticProperty("FailureCount");
    Assert.AreEqual(failureCount, 0);
}
```

Again, you create an instance of `PrivateType`, passing the type reference for the class you wish to access. Then you use that instance, invoking `GetStaticProperty` to retrieve the value you wish to test. Finally, you ensure that the value is zero as expected.

Use caution when testing static data. Because static data is shared and is not automatically reset between your tests, sometimes the order of your tests will affect their results. In other words, if you test that a value is initialized to zero in one test and then set it to a test value in another test, if the order of those tests is reversed, the zero test will fail. Remember that you must be able to run your tests in any order and in any combination.

Code Generation

Remember the work you did earlier in this chapter to create your first unit test? Depending upon the degree to which you practice TDD, you might not create many of your unit tests in that manner. We created the first unit tests manually to help convey basic concepts, but Team System has support for automatically generating code. You may generate unit tests from your implementation code or generate limited implementation code when writing your tests.

Generating tests from code

If you have ever needed to add unit testing to an existing project that had none, it was likely a frustrating experience. Fortunately, Team System has introduced the capability to generate outlines of unit tests based on selected implementation code.

If you are practicing test-driven development, described earlier, be aware that generating tests from existing code is considered a very non-TDD practice. In pure TDD, no implementation code should exist before unit tests are created.

Let's begin with the `Functions` class we used earlier in this chapter. Open `Functions.cs` and ensure it contains the following code:

```
using System;

namespace ExtendedMath
{
    public sealed class Functions
    {
```

```
private Functions() {}

public static int Fibonacci(int factor)
{
    if (factor < 2)
        return (factor);
    int x = Fibonacci(--factor);
    int y = Fibonacci(--factor);

    return x + y;
}
}
```

If you have been following the examples in this chapter, delete your existing `FunctionsTest.cs` file. Now you can right-click in your code and choose `Create Unit Tests` or choose the `Create Tests` button from the `Test Views` toolbar. The `Create Unit Tests` dialog will appear, as shown in Figure 14-7.




Figure 14-7

Right-clicking in code is context-sensitive and the dialog will default appropriately based on where you clicked. For example, if you click from your class definition, all of that class's members will be selected for generation by default. Clicking on a method will default with only that method selected.

Select the members for which you would like to generate unit tests and then click the `Settings` button. The `Test Generation Settings` dialog will appear, as shown in Figure 14-8.

This dialog enables you to choose the behavior and naming conventions that will be used to generate your unit test code. For this first run, leave the settings at their defaults and click `OK`.

Team System will create a new class file, `FunctionsTest.cs`, if it does not already exist. Inside, you will find a class, `FunctionsTest`, marked with `[TestClass]` attribute.

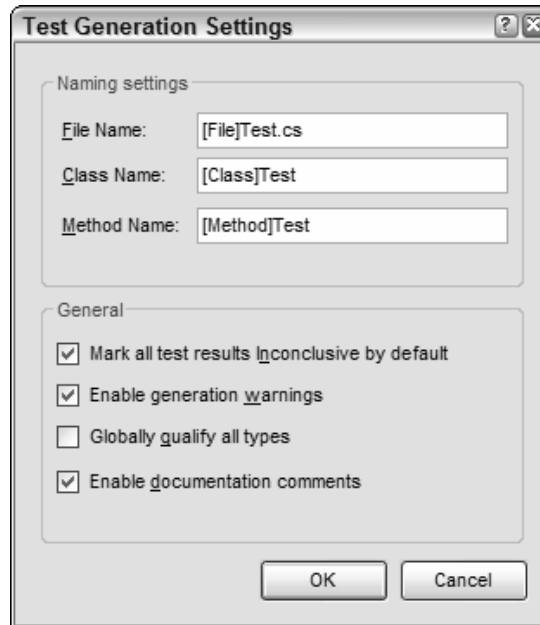


Figure 14-8

At the top of the class is a `TestContext` field initialized to `null`, with a `TestContext` property to enable you to access context from your tests:

```
private TestContext testContextInstance;

public TestContext TestContext
{
    get
    {
        return testContextInstance;
    }
    set
    {
        testContextInstance = value;
    }
}
```

Next, you will find commented-out placeholder methods for `ClassInitialize`, `ClassCleanup`, `TestInitialize`, and `TestCleanup`, wrapped in a region:

```
#region Additional test attributes
//
//You can use the following additional attributes as you write your tests:
//
//Use ClassInitialize to run code before running the first test in the
class
//
```

```
//[ClassInitialize()]
//public static void MyClassInitialize(TestContext testContext)
//{
//}
//
//Use ClassCleanup to run code after all tests in a class have run
//
//[ClassCleanup()]
//public static void MyClassCleanup()
//{
//}
//
//Use TestInitialize to run code before running each test
//
//[TestInitialize()]
//public void MyTestInitialize()
//{
//}
//
//Use TestCleanup to run code after each test has run
//
//[TestCleanup()]
//public void MyTestCleanup()
//{
//}
//
#endregion
```

Finally, you will see the actual generated unit test:

```
/// <summary>
/// A test for Fibonacci (int)
/// </summary>
[TestMethod()]
public void FibonacciTest()
{
    int factor = 0; // TODO: Initialize to an appropriate value

    int expected = 0;
    int actual;

    actual = ExtendedMath.Functions.Fibonacci(factor);

    Assert.AreEqual(expected, actual, "ExtendedMath.Functions.Fibonacci did
not return the expected value.");
    Assert.Inconclusive("Verify the correctness of this test method.");
}
```

The generated code defines a basic structure that depends on the signature of the member being tested. In this example, it recognized that `Fibonacci` accepts an integer and returns an integer, including an `Assert.AreEqual` for you. The `TODO` indicates that `factor` and `expected` are only default values and need to be adjusted by you to ensure correct test values.

Chapter 14

Keep in mind that generated code will often benefit from careful refactoring, and your generated unit tests will be no exception. For example, look for ways to consolidate common tasks into shared functions. Refactoring is covered in Chapter 11.

Optionally, but by default, generated tests end with calls to `Assert.Inconclusive` to indicate that the unit test needs further inspection to determine its correctness. See the “Using the Assert methods” section for more information.

Generating code from tests

As you have just seen, generating tests from existing code can be very useful. It can be invaluable when you’ve taken ownership of a system that has no unit tests. However, adherents of test-driven development discourage such approaches, preferring to write test code before implementation code.

Fortunately, Visual Studio also has some support for this approach. You can write unit tests, including calls to classes and methods that do not exist, and the IDE will help you define outlines of their implementations automatically.

Begin by creating a new Class Library project called `ResourceExample`. While TDD purists prefer to avoid writing any implementation code before tests, Team System requires an existing class in which to place generated code. All you need is a basic class definition, so add a new class file called `ResourceLoader.cs` containing the following code:

```
using System;

namespace ResourceExample
{
    public class ResourceLoader
    {
    }
}
```

Now you can write your unit test. It will verify that a call to the `ResourceLoader` class’s `GetText` method with an identifier returns the expected resource string. Create a new test project named `ResourceExampleTesting` and set a reference to the `ResourceExample` project. Next, add a new class file called `ResourceLoaderTest.cs` with the following code:

```
using Microsoft.VisualStudio.TestTools.UnitTesting;
using ResourceExample;

namespace ResourceExampleTesting
{
    [TestClass]
    public class ResourceLoaderTest
    {
        [TestMethod]
        public void GetTextTest()
        {
            const int RESOURCE = 100;
            const string EXPECTED = "Result for 100";
            string actual = ResourceLoader.GetText(RESOURCE);

            Assert.AreEqual(EXPECTED, actual);
        }
    }
}
```

```
    }  
  }  
}
```

In the code editor, click on `GetText`. It should be underscored with a small rectangle indicating that one or more options are available. Hover the mouse pointer over the rectangle, click on the resulting button, and you will see “Generate method stub for ‘`GetText`’ in ‘`ResourceExample.ResourceLoader`.’”

Alternately, you can right-click on `GetText` and choose `Generate Method Stub`. *If you have trouble finding these options, make sure that you have set a reference from the test project to the `ResourceExample` project.*

Choose either of these options and then switch to `ResourceLoader.cs`, where you will now see the following:

```
using System;  
  
namespace ResourceExample  
{  
    public class ResourceLoader  
    {  
        public static string GetText(int RESOURCE)  
        {  
            throw new Exception("The method or operation is not implemented.");  
        }  
    }  
}
```

You will notice that Visual Studio inferred the best signature for the method from the context of its use in the unit test. Also notice that the only action in the method is throwing an `Exception`. This is to remind you that there is work left to do. Replace the `throw` with whatever implementation code you need to make your test(s) pass, and then refactor and continue with your next test.

Unit Testing ASP.NET Applications

A common problem for developers prior to Team System was the challenge of unit testing ASP.NET applications. Most of the behavior of ASP.NET applications relies on having a web request context—for example, querystrings, cookies, and server variables. To effectively test this code, a web environment needed to be simulated. While some frameworks, such as `NUnitASP` existed, the task was still difficult and not integrated with the Visual Studio environment.

Team System enables developers to write unit tests that have access to an active ASP.NET context. The unit testing framework connects to the target ASP.NET application, executes a request, and provides the context of that request to the unit test.

ASP.NET unit test attributes

An ASP.NET unit test looks much like any other unit test, except that it is decorated with additional attributes to indicate details about the target ASP.NET application. The following table summarizes common ASP.NET unit testing attributes.

Attribute	Description
HostType	Identifies the type of host that will run the unit test. Typically, this will be ASP.NET.
UrlToTest	The URL that will be requested. The context of this request will be made available to the unit test via <code>TestContext</code> .
AspNetDevelopmentServerHost	Specify this attribute when the host for the web application will be the built-in ASP.NET Development Server. The first argument is the physical directory of the web application. The second is the name of the application root.

These attributes exist in the `Microsoft.VisualStudio.TestTools.UnitTesting.Web` namespace, and, as with all unit tests, a reference to the `Microsoft.VisualStudio.TestTools.UnitTestingFramework` assembly is required. Ensure that you add the associated using statements as well:

```
using Microsoft.VisualStudio.TestTools.UnitTesting;  
using Microsoft.VisualStudio.TestTools.UnitTesting.Web;
```

Now that you have an understanding of the attributes for ASP.NET unit tests, we'll describe your options for creating them.

Creating ASP.NET unit tests

One way to create an ASP.NET unit test is to generate one from a class or method of your web application. You saw in the previous section how simple Team System makes this.

First, open your web application project. Then, open the class containing methods you wish to test. Right-click and choose **Create Unit Tests**. Select the specific methods for which you wish to generate unit tests and click **OK**.

The unit tests, already decorated with the appropriate attributes for your web application, are created in a test project in your solution. If you don't already have a test project, one will be created automatically for you.

ASP.NET unit test generation is only available to classes in your web application's `App_Code` folder. For other items, such as code-behind implementations, we suggest you copy the attributes from a previously generated ASP.NET unit test to another unit test for the code you wish to test.

Whether you generated your ASP.NET unit test or converted a unit test by adding attributes, you are ready to tap into the ASP.NET context in order to verify behavior. The `TestContext` object features a `RequestedPage` property. This is actually an instance of a `System.Web.UI.Page` object appropriate for your web application. Assign this to a local variable and you can use it to access controls of the page.

In order to read and set values of the controls on the page, you will need to obtain a reference to the control using the `Page` object's `FindControl` method.

Frequently, you will need to access nonpublic methods of the page. For example, most button submit methods are not public. As you saw earlier in this chapter, Team System provides a simple way to wrap objects in order to access their nonpublic members. To access such methods, use a `PrivateObject` pointing to the `Page` instance. Once you have that `PrivateObject`, use its `Invoke` method to call the target nonpublic method.

Here is an example ASP.NET unit test, using a local ASP.NET development server, which enters a stock ticker symbol and requests a quote:

```
[TestMethod]
[HostType("ASP.NET")]
[AspNetDevelopmentServerHost("%PathToWebRoot%\StockTicker", "/StockTicker")]
[UrlToTest("http://localhost/StockTicker")]
public void VerifyCompanyTickerLookup()
{
    Page page = TestContext.RequestedPage;
    TextBox stockName = (TextBox)page.FindControl("txtStockSymbol");
    stockName.Text = "MSFT";

    PrivateObject po = new PrivateObject(page);
    Button getCompanyName = (Button)page.FindControl("cmdGetCompanyName");
    po.Invoke("cmdGetCompanyName_Click", getCompanyName, EventArgs.Empty);

    Label result = (Label)page.FindControl("lblResults");
    Assert.AreEqual("Microsoft", result.Text);
}
```

The attributes indicate how to connect to the local `StockTicker` ASP.NET web application. First, the `Page` reference is obtained from the `TestContext`. Then, the value of the `txtStockName` control is set to `MSFT`. Access to the `Submit` button is not public, so a `PrivateObject` is used to access the `Page`. After invoking the button's click method, the `txtResult` control is read to ensure that the ticker `MSFT` was correctly returned as `"Microsoft"`.

Code Coverage

The unit testing features in Team System have full support for code coverage. Code coverage automatically inserts tracking logic, a process called *instrumentation*, to monitor which lines of code are executed during the execution of your tests. The most important result of this is the identification of regions of your code that you have not reached with your tests.

Often, you may have branching or exception-handling logic that isn't executed in common situations. It is critical to use code coverage to identify these areas because your users certainly will. Add unit tests to cause those areas of code to be executed, and you'll be able to sleep soundly at night.

Code coverage is a useful tool, but it should not be relied upon as an exclusive indicator of unit test effectiveness. It cannot tell you the manner in which code was executed, possibly missing errors that would result with different data or timing. A suite of unit tests based on a variety of different inputs and execution orders will help to ensure that your code is correct, complete, and resilient. Use code coverage to help identify code your tests are missing, not to tell you when your tests are complete.

Chapter 14

A tenet of effective unit testing is that the removal of any line of code should cause at least one unit test to fail. This is, of course, an ideal, but worth keeping in mind as you develop your systems.

Enabling code coverage

Code coverage is activated via a setting in the Test Run Configuration. Open the configuration for editing by choosing Test ⇄ Edit Test Run Configuration and selecting the configuration. Once you have the Edit Test Run Configuration dialog active, select the Code Coverage page.

Select the assemblies you wish to instrument from the Select Artifacts to Instrument list. If you don't see an assembly you'd like to instrument, you can click Add Assembly to manually add it. Figure 14-9 is a screenshot of the Code Coverage page of the Test Run Configuration editor.

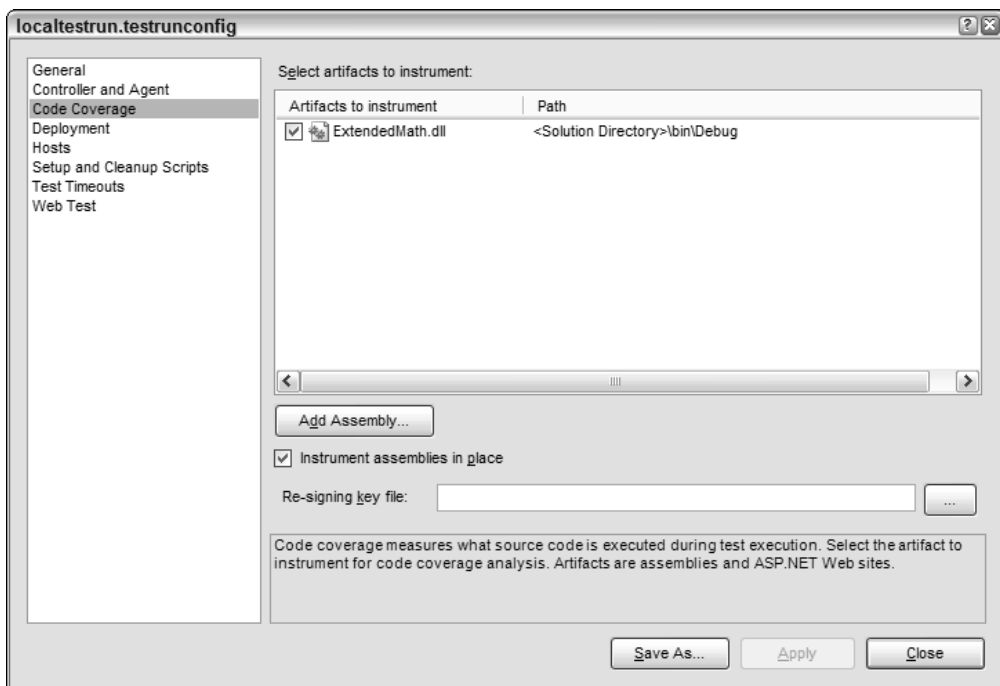


Figure 14-9

The instrumentation process modifies your original assemblies, invalidating original signatures. If you are working with signed assemblies, use the Re-signing key file field to specify a key file with which to sign the instrumented assembly.

Viewing code coverage results

Once you have enabled coverage and selected the assemblies to instrument, run your unit tests as normal. You will then be able to see results of the coverage in the Code Coverage Results window. This window will show counts of lines and percentages of code covered and uncovered. You may expand the view by clicking the plus signs to see details at the assembly, class, and member levels.

Unit Testing with the Unit Test Framework

To quickly determine which areas of code need attention, enable code coverage highlighting of your code by pressing the Show Code Coverage button on the Code Coverage toolbar. Executable lines of code will be highlighted in red if they have not been run by your tests and in blue if they were. Code that is purely structural or documentation will not be highlighted. Figure 14-10 illustrates the results of a code coverage test.

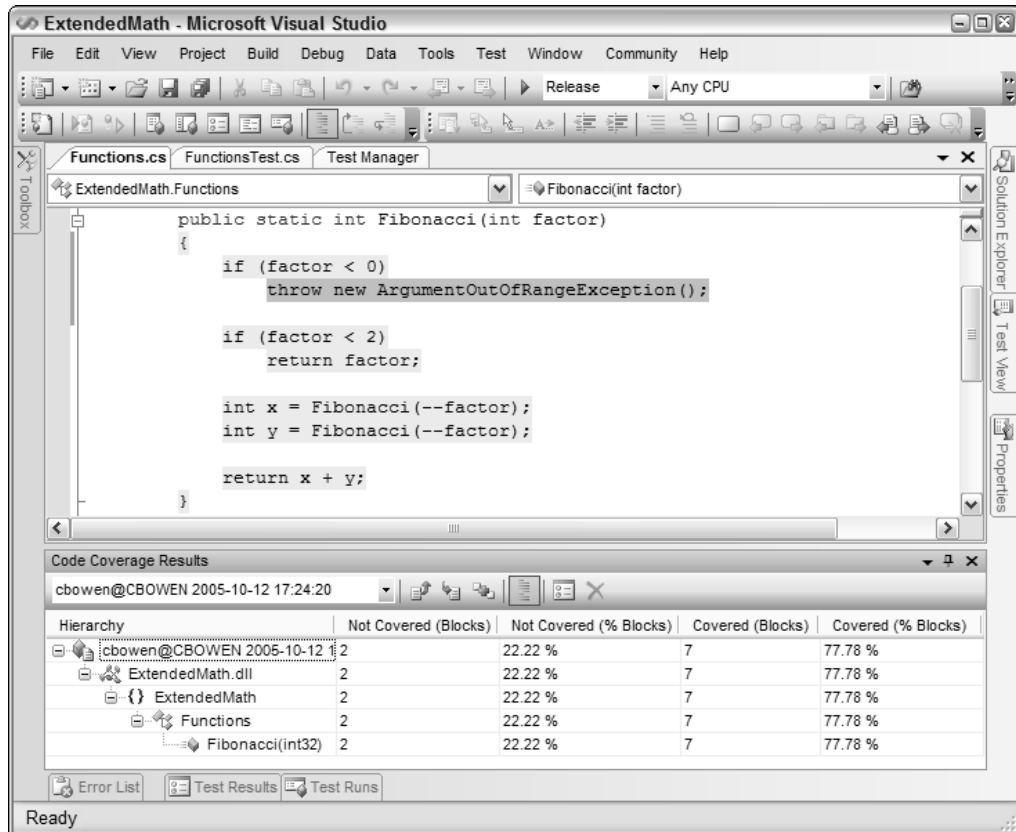


Figure 14-10

In this example, we added two lines to our previous `Fibonacci` implementation that check for a negative factor and throw an exception. Code Coverage has colored the line throwing the exception red. This means that none of the unit tests we executed caused that line to run. This is a clear indication that we need another unit test. Create a new unit test—for example, `FibonacciOfNegativeFactorsNotSupportedTest`—and test calling `Fibonacci` with a factor less than zero. Decorate the test with the `ExpectedException` attribute, indicating that you expect the `ArgumentOutOfRangeException` in order for the test to pass. Rerun the unit tests and the `Fibonacci` method should now have 100 percent coverage.

Again, keep in mind that 100 percent code coverage does not mean you have finished writing your unit tests. Proper testing may involve multiple executions of the same code using different data. Code coverage is one measure of effectiveness, but certainly not the only one. Consider adopting test-driven development, plan your testing cases up front, and then use code coverage to alert you to scenarios you forgot to test.

Summary

The Developer and Tester editions of Visual Studio Team System bring the advantages of unit testing to the developer by fully integrating features with the development environment. If you're new to unit testing, this chapter has provided an overview of what unit testing is, how to create effective unit tests, and some of the principles behind test-driven development. We covered the creation and management of unit tests and detailed the methods and attributes available in the unit test framework. You should be familiar with attributes for identifying your tests as well as many of the options that the `Assert` class offers for testing behavior of code.

You've learned how to generate code with Team System, either by generating tests from code or code from tests. As you saw, it is also very easy to bind unit tests to a data source to create data-driven tests. This chapter also showed that Team System offers a simple way to access private members for testing, addressing a need that previously required manual workarounds.

You've also seen that ASP.NET developers can now create unit tests against web applications without relying on external tools and frameworks. Finally, you learned how to use code coverage to help identify where your unit tests may have missed some scenarios.

If you have been using other frameworks for your unit testing to date, you'll enjoy not having to leave the IDE to work with your tests, as well as the familiar options and syntax that Team System offers.

We strongly encourage you to become familiar with the benefits of unit testing, keeping in mind that unit tests are not a replacement for other forms of testing, but a very strong supplement.

Over the next three chapters, we'll continue our look at those other forms of testing that Team System Team Tester Edition supports. You'll learn about load and web tests in Chapter 15, manual tests in Chapter 16, and custom tests in Chapter 17.